

Paging in RISC-V (SV39)

In this chapter, we introduce **virtual memory** to Novix OS by implementing the **SV39 paging system** — the 39-bit virtual memory scheme used by RISC-V 64-bit systems.

Paging provides **isolation**, **protection**, and **address space translation** for every process.

Concept Overview

In **SV39**, each virtual address (VA) is divided into:

- 9 bits VPN[2]
- 9 bits VPN[1]
- 9 bits VPN[0]
- 12 bits page offset

This structure creates a **3-level page table**:

- **Level 2 (Root)** → Page directory pointer
- **Level 1 (Middle)** → Page directory
- **Level 0 (Leaf)** → Page table entries (PTEs)

Each PTE describes the mapping between a **virtual page** and a **physical frame** and includes permission bits:

| Flag | Meaning |
|------|-----------------|
| V | Valid entry |
| R | Read |
| W | Write |
| X | Execute |
| U | User accessible |
| G | Global |
| A | Accessed |
| D | Dirty |

Function: `paging_init`

```
void paging_init(void);
```

Purpose

Initializes the kernel's root page table and performs **identity mapping** of the kernel memory.

Detailed Steps

1. Allocate one page for the root page table:

```
paddr_t pt_page = alloc_pages(1);
kernel_pagetable = (pagetable_t)pt_page;
```

2. Identity-map all kernel memory from `_kernel_base` to `_heap_end`:

```
map_page(kernel_pagetable, pa, pa, PTE_V | PTE_R | PTE_W | PTE_X, 0);
```

3. Map MMIO regions for devices such as **VirtIO**:

```
map_mmio_range(kernel_pagetable, VIRTIO_MMIO_BASE, VIRTIO_MMIO_BASE,
                VIRTIO_MMIO_SIZE * VIRTIO_MMIO_MAX_DEVICES,
                PTE_V | PTE_R | PTE_W);
```

- Load this new page table into the CPU:

```
set_active_pagetable((uintptr_t)kernel_pagetable);
```

- Log success:

```
uart_printf("[paging_init] kernel page table initialized at : 0x%lx\n",
            (uintptr_t)kernel_pagetable);
```

Function: alloc_pages

```
paddr_t alloc_pages(uint64_t n);
```

Purpose

Allocates one or more contiguous physical pages.

This is the physical memory backend used by the paging subsystem.

Algorithm

- For a single page:
 - Try to reuse one from the **free list** (`alloc_free_list()`)
 - Otherwise use the **bump allocator** (`next_paddr`)
- For multiple pages:
 - Allocate sequentially via bump pointer
 - Zero-initialize using `memset()`
 - Track each allocation with `track_alloc()`

Function: free_page / free_pages_range

```
void free_page(paddr_t pa);
void free_pages_range(paddr_t pa, size_t npages);
```

Purpose

Releases one or more physical pages and returns them to the free list.

Description

- Adds the page's address to the free list.
- Increments `g_free_pages` counter.
- Tracks the deallocation for debugging purposes (`track_free()`).

Function: map_page

```
void map_page(pagetable_t root_table, uint64_t va, uint64_t pa,
              uint64_t flags, int debug_flag);
```

Purpose

Creates a mapping from **virtual address (VA)** to **physical address (PA)**.

Detailed Description

1. Extract indices:

```
vpn2 = (va >> 30) & 0x1FF;  
vpn1 = (va >> 21) & 0x1FF;  
vpn0 = (va >> 12) & 0x1FF;
```

2. Traverse or allocate new page tables at each level (L2 → L1 → L0).

3. Insert the final **leaf PTE** at level 0:

```
pt[vpn0] = ((pa >> 12) << 10) | flags | PTE_V;
```

4. Optional UART log when `debug_flag` is true.

Function: `unmap_page`

```
void unmap_page(pagetable_t root_table, vaddr_t va);
```

Purpose

Removes a mapping from a virtual address.

Behavior

- Traverses the same 3-level page table.
- Clears the leaf PTE.
- Executes `sfence_vma()` to flush the TLB.

Function: `set_active_pagetable`

```
void set_active_pagetable(uintptr_t new_pagetable);
```

Purpose

Switches the CPU's currently active page table.

Mechanism

- Writes the `satp` CSR with the new page table physical address:

```
WRITE_CSR(satp, MAKE_SATP(new_pagetable));  
sfence_vma();
```

Function: `map_mmio_range`

```
void map_mmio_range(pagetable_t pt, uint64_t pa_start,  
                     uint64_t va_start, uint64_t len, uint64_t perm);
```

Purpose

Maps **memory-mapped I/O (MMIO)** regions such as device buffers.

Implementation

Loops over each page of the MMIO region and calls `map_page()` with the specified permissions.

Function: walk_page

```
paddr_t walk_page(pagetable_t pagetable, vaddr_t va);
```

Purpose

Traverses a page table and returns the physical address mapped to a virtual address.

Return Value

- Physical address (if mapping exists)
- 0 if unmapped or invalid

Function: walk

```
pte_t* walk(pagetable_t pagetable, uint64_t va, int alloc);
```

Purpose

Locates or creates the PTE for a given virtual address.

If `alloc` is true, missing page tables are allocated dynamically.

Function: walkaddr

```
paddr_t walkaddr(pagetable_t pagetable, vaddr_t va);
```

Purpose

Translates a user-space virtual address to a physical address.

Safety

- Rejects kernel-space addresses.
- Verifies the PTE's validity and permission bits.
- Returns 0 on failure.

Function: free_pagetable

```
void free_pagetable(pagetable_t pt, int lvl);
```

Purpose

Recursively frees all page-table pages belonging to a process.

Behavior

- Traverses all 512 entries per level.
- Frees only **user pages** (PTE_U set).
- Skips kernel mappings.
- Clears PTE entries after freeing.

Function: debug_pagetable

```
void debug_pagetable(const char *who);
```

Purpose

Prints the currently active page table's physical address by reading the `satp` CSR.

Function: `page_allocator_free_pages`

```
uint64_t page_allocator_free_pages(void);
```

Purpose

Returns the count of currently free physical pages (`g_free_pages`).

Function: `kernel_main`

```
void kernel_main(void);
```

Purpose

Initializes all major kernel subsystems, including paging.

Steps

1. Clear BSS

```
bss_init();
```

2. Boot message

```
uart_cls();
uart_printf("Novix RISC-V 64 OS, (c) NovixManiac, Version 0.0.1\n\n");
```

3. Initialize subsystems

```
trap_init();
stack_init();
ram_init();
heap_init();
```

4. Enable Paging

```
paging_init();
```

5. Enter infinite loop.

Summary

At this stage, **Novix OS** now has a functioning **virtual memory system** powered by the **SV39** paging scheme. You have implemented:

- Multi-level page tables (L2 → L1 → L0)
- Physical page allocation and freeing
- Page mapping, unmapping, and lookup
- Kernel identity mapping
- MMIO region handling
- Page-table switching via `satp`

This is one of the most important architectural milestones — from now on, every address in Novix OS goes through the paging layer, providing the basis for **process isolation**, **system calls**, and **virtual address spaces**.

© NovixManiac — *The Art of Operating Systems Development*